

ML Task-01

Implement a linear regression model to predict the prices of houses based on their square footage and the number of bedrooms and bathrooms.

Dataset : - <https://www.kaggle.com/c/house-prices-advanced-regression-techniques/data>

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

df_train = pd.read_csv('train.csv')
df_test = pd.read_csv('test.csv')

df_train.head()
```

| | Id | MSSubClass | MSZoning | LotFrontage | LotArea | Street | Alley | LotShape |
|---|----|------------|----------|-------------|---------|--------|-------|----------|
| 0 | 1 | 60 | RL | 65.0 | 8450 | Pave | NaN | Reg |
| 1 | 2 | 20 | RL | 80.0 | 9600 | Pave | NaN | Reg |
| 2 | 3 | 60 | RL | 68.0 | 11250 | Pave | NaN | IR1 |
| 3 | 4 | 70 | RL | 60.0 | 9550 | Pave | NaN | IR1 |
| 4 | 5 | 60 | RL | 84.0 | 14260 | Pave | NaN | IR1 |

| | LandContour | Utilities | ... | PoolArea | PoolQC | Fence | MiscFeature | MiscVal |
|----|-------------|-----------|-----|----------|--------|-------|-------------|---------|
| 0 | Lvl | AllPub | ... | 0 | NaN | NaN | NaN | 0 |
| 2 | Lvl | AllPub | ... | 0 | NaN | NaN | NaN | 0 |
| 5 | Lvl | AllPub | ... | 0 | NaN | NaN | NaN | 0 |
| 9 | Lvl | AllPub | ... | 0 | NaN | NaN | NaN | 0 |
| 12 | Lvl | AllPub | ... | 0 | NaN | NaN | NaN | 0 |

| | YrSold | SaleType | SaleCondition | SalePrice |
|---|--------|----------|---------------|-----------|
| 0 | 2008 | WD | Normal | 208500 |
| 1 | 2007 | WD | Normal | 181500 |
| 2 | 2008 | WD | Normal | 223500 |
| 3 | 2006 | WD | Abnorml | 140000 |

```
4    2008         WD         Normal    250000
```

```
[5 rows x 81 columns]
```

```
df_test.head()
```

| | Id | MSSubClass | MSZoning | LotFrontage | LotArea | Street | Alley |
|---|------|------------|----------|-------------|---------|--------|-------|
| 0 | 1461 | 20 | RH | 80.0 | 11622 | Pave | NaN |
| 1 | 1462 | 20 | RL | 81.0 | 14267 | Pave | NaN |
| 2 | 1463 | 60 | RL | 74.0 | 13830 | Pave | NaN |
| 3 | 1464 | 60 | RL | 78.0 | 9978 | Pave | NaN |
| 4 | 1465 | 120 | RL | 43.0 | 5005 | Pave | NaN |


| | LandContour | Utilities | ... | ScreenPorch | PoolArea | PoolQC | Fence |
|---|-------------|-----------|-----|-------------|----------|--------|-------|
| 0 | Lvl | AllPub | ... | 120 | 0 | NaN | MnPrv |
| 1 | Lvl | AllPub | ... | 0 | 0 | NaN | NaN |
| 2 | Lvl | AllPub | ... | 0 | 0 | NaN | MnPrv |
| 3 | Lvl | AllPub | ... | 0 | 0 | NaN | NaN |
| 4 | HLS | AllPub | ... | 144 | 0 | NaN | NaN |

| | MiscVal | MoSold | YrSold | SaleType | SaleCondition |
|---|---------|--------|--------|----------|---------------|
| 0 | 0 | 6 | 2010 | WD | Normal |
| 1 | 12500 | 6 | 2010 | WD | Normal |
| 2 | 0 | 3 | 2010 | WD | Normal |
| 3 | 0 | 6 | 2010 | WD | Normal |
| 4 | 0 | 1 | 2010 | WD | Normal |

```
[5 rows x 80 columns]
```

```
sns.heatmap(df_train.isnull())
```

```
<Axes: >
```



```
df_train.columns
Index(['Id', 'MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea',
       'Street',
       'Alley', 'LotShape', 'LandContour', 'Utilities', 'LotConfig',
       'LandSlope', 'Neighborhood', 'Condition1', 'Condition2',
       'BldgType',
       'HouseStyle', 'OverallQual', 'OverallCond', 'YearBuilt',
       'YearRemodAdd',
       'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd',
       'MasVnrType',
       'MasVnrArea', 'ExterQual', 'ExterCond', 'Foundation',
       'BsmtQual',
       'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1',
       'BsmtFinType2', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF',
       'Heating',
       'HeatingQC', 'CentralAir', 'Electrical', '1stFlrSF',
       '2ndFlrSF',
       'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath',
       'FullBath',
```

```

        'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual',
        'TotRmsAbvGrd', 'Functional', 'Fireplaces', 'FireplaceQu',
        'GarageType',
        'GarageYrBlt', 'GarageFinish', 'GarageCars', 'GarageArea',
        'GarageQual',
        'GarageCond', 'PavedDrive', 'WoodDeckSF', 'OpenPorchSF',
        'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea',
        'PoolQC',
        'Fence', 'MiscFeature', 'MiscVal', 'MoSold', 'YrSold',
        'SaleType',
        'SaleCondition', 'SalePrice'],
        dtype='object')

```

```

null_columns=[]
for i in df_train.columns.tolist():
    if df_train[i].isnull().sum() >= 500:
        null_columns.append(i)
        print(i,df_train[i].isnull().sum())

```

```

Alley 1369
MasVnrType 872
FireplaceQu 690
PoolQC 1453
Fence 1179
MiscFeature 1406

```

```

df_train = df_train.drop(columns=null_columns)
df_test = df_test.drop(columns=null_columns)

```

```

df_train_cleaned = df_train.ffill()
df_test_cleaned = df_test.ffill()

```

```

df_train_cleaned.head()

```

| | Id | MSSubClass | MSZoning | LotFrontage | LotArea | Street | LotShape |
|---------------|--------|------------|----------|-------------|---------|--------|----------|
| LandContour \ | | | | | | | |
| 0 | 1 | 60 | RL | 65.0 | 8450 | Pave | Reg |
| Lvl | | | | | | | |
| 1 | 2 | 20 | RL | 80.0 | 9600 | Pave | Reg |
| Lvl | | | | | | | |
| 2 | 3 | 60 | RL | 68.0 | 11250 | Pave | IR1 |
| Lvl | | | | | | | |
| 3 | 4 | 70 | RL | 60.0 | 9550 | Pave | IR1 |
| Lvl | | | | | | | |
| 4 | 5 | 60 | RL | 84.0 | 14260 | Pave | IR1 |
| Lvl | | | | | | | |
| Utilities | | | | | | | |
| LotConfig \ | | | | | | | |
| PoolArea \ | | | | | | | |
| 0 | AllPub | Inside | ... | 0 | 0 | 0 | |
| 0 | | | | | | | |


| | | | | | | |
|---|--------|--------|-----|-----|---|---|
| 1 | AllPub | FR2 | ... | 0 | 0 | 0 |
| 0 | | | | | | |
| 2 | AllPub | Inside | ... | 0 | 0 | 0 |
| 0 | | | | | | |
| 3 | AllPub | Corner | ... | 272 | 0 | 0 |
| 0 | | | | | | |
| 4 | AllPub | FR2 | ... | 0 | 0 | 0 |
| 0 | | | | | | |

| | MiscVal | MoSold | YrSold | SaleType | SaleCondition | SalePrice |
|---|---------|--------|--------|----------|---------------|-----------|
| 0 | 0 | 2 | 2008 | WD | Normal | 208500 |
| 1 | 0 | 5 | 2007 | WD | Normal | 181500 |
| 2 | 0 | 9 | 2008 | WD | Normal | 223500 |
| 3 | 0 | 2 | 2006 | WD | Abnorml | 140000 |
| 4 | 0 | 12 | 2008 | WD | Normal | 250000 |

[5 rows x 75 columns]

```
sns.heatmap(df_train_cleaned.isnull())
```

<Axes: >



```
df_train_cleaned.describe()
```

| | Id | MSSubClass | LotFrontage | LotArea |
|---------------|-------------|-------------|-------------|---------------|
| OverallQual \ | | | | |
| count | 1460.000000 | 1460.000000 | 1460.000000 | 1460.000000 |
| mean | 730.500000 | 56.897260 | 70.104795 | 10516.828082 |
| std | 421.610009 | 42.300571 | 23.846996 | 9981.264932 |
| min | 1.000000 | 20.000000 | 21.000000 | 1300.000000 |
| 25% | 365.750000 | 20.000000 | 59.000000 | 7553.500000 |
| 50% | 730.500000 | 50.000000 | 70.000000 | 9478.500000 |
| 75% | 1095.250000 | 70.000000 | 80.000000 | 11601.500000 |
| max | 1460.000000 | 190.000000 | 313.000000 | 215245.000000 |

| | OverallCond | YearBuilt | YearRemodAdd | MasVnrArea |
|------------------|-------------|-------------|--------------|-------------|
| BsmtFinSF1 ... \ | | | | |
| count | 1460.000000 | 1460.000000 | 1460.000000 | 1460.000000 |
| 1460.000000 ... | | | | |
| mean | 5.575342 | 1971.267808 | 1984.865753 | 103.492466 |
| 443.639726 ... | | | | |
| std | 1.112799 | 30.202904 | 20.645407 | 180.795612 |
| 456.098091 ... | | | | |
| min | 1.000000 | 1872.000000 | 1950.000000 | 0.000000 |
| 0.000000 ... | | | | |
| 25% | 5.000000 | 1954.000000 | 1967.000000 | 0.000000 |
| 0.000000 ... | | | | |
| 50% | 5.000000 | 1973.000000 | 1994.000000 | 0.000000 |
| 383.500000 ... | | | | |
| 75% | 6.000000 | 2000.000000 | 2004.000000 | 165.250000 |
| 712.250000 ... | | | | |
| max | 9.000000 | 2010.000000 | 2010.000000 | 1600.000000 |
| 5644.000000 ... | | | | |

| | WoodDeckSF | OpenPorchSF | EnclosedPorch | 3SsnPorch |
|---------------|-------------|-------------|---------------|-------------|
| ScreenPorch \ | | | | |
| count | 1460.000000 | 1460.000000 | 1460.000000 | 1460.000000 |
| 1460.000000 | | | | |
| mean | 94.244521 | 46.660274 | 21.954110 | 3.409589 |
| 15.060959 | | | | |
| std | 125.338794 | 66.256028 | 61.119149 | 29.317331 |
| 55.757415 | | | | |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 0.000000 | | | | |
| 25% | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 0.000000 | | | | |
| 50% | 0.000000 | 25.000000 | 0.000000 | 0.000000 |
| 0.000000 | | | | |
| 75% | 168.000000 | 68.000000 | 0.000000 | 0.000000 |
| 0.000000 | | | | |
| max | 857.000000 | 547.000000 | 552.000000 | 508.000000 |
| 480.000000 | | | | |

| | PoolArea | MiscVal | MoSold | YrSold |
|---------------|-------------|-------------|-------------|-------------|
| SalePrice | | | | |
| count | 1460.000000 | 1460.000000 | 1460.000000 | 1460.000000 |
| 1460.000000 | | | | |
| mean | 2.758904 | 43.489041 | 6.321918 | 2007.815753 |
| 180921.195890 | | | | |
| std | 40.177307 | 496.123024 | 2.703626 | 1.328095 |
| 79442.502883 | | | | |
| min | 0.000000 | 0.000000 | 1.000000 | 2006.000000 |
| 34900.000000 | | | | |
| 25% | 0.000000 | 0.000000 | 5.000000 | 2007.000000 |

```

129975.000000
50%      0.000000      0.000000      6.000000  2008.000000
163000.000000
75%      0.000000      0.000000      8.000000  2009.000000
214000.000000
max      738.000000  15500.000000  12.000000  2010.000000
755000.000000


```

```
[8 rows x 38 columns]
```

```

plt.title('Salary Distribution Plot')
sns.histplot(df_train_cleaned['SalePrice'])
plt.show()

```



Split data

```

from sklearn.model_selection import train_test_split

features = ['GrLivArea', '1stFlrSF', '2ndFlrSF', 'LowQualFinSF',
            'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath',
            'BedroomAbvGr']

```



```


target = 'SalePrice'
X = df_train_cleaned[features]
y = df_train_cleaned[target]

df = pd.concat([X, y], axis=1)

# Plotting grid of scatter plots for each feature against the target
variable
sns.pairplot(df, x_vars=features, y_vars=[target], kind='scatter',
              plot_kws={'alpha':0.4}, diag_kws={'alpha':0.55,
              'bins':40})

plt.show()

```




```


sns.lmplot(x='GrLivArea',
            y='SalePrice',
            data=df_train_cleaned,
            scatter_kws={'alpha':0.3})

<seaborn.axisgrid.FacetGrid at 0x17727b890>


```



```
sns.lmplot(x='BedroomAbvGr',  
           y='SalePrice',  
           data=df_train_cleaned,  
           scatter_kws={'alpha':0.3})  
<seaborn.axisgrid.FacetGrid at 0x177079670>
```



```
sns.lmplot(x='FullBath',  
           y='SalePrice',  
           data=df_train_cleaned,  
           scatter_kws={'alpha':0.3})  
  
<seaborn.axisgrid.FacetGrid at 0x176f73ec0>
```



```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =  
0.2, random_state = 42)
```

Train model (Linear Regression)


```
from sklearn.linear_model import LinearRegression  
  
model = LinearRegression()  
model.fit(X_train, y_train)  
  
LinearRegression()  
coefficients = model.coef_  
model.score(X, y)  
0.6541686106425739
```

```
for feature, coef in zip(features, coefficients):  
    print(f'{feature}: {coef}')
```

```
GrLivArea: 45.45358200866404  
1stFlrSF: 70.93123244880826  
2ndFlrSF: 19.004245633922277  
LowQualFinSF: -44.481896073804585  
BsmtFullBath: 22125.53078348177  
BsmtHalfBath: 6538.890647874652  
FullBath: 34869.22887902238  
HalfBath: 22717.7813535538  
BedroomAbvGr: -18379.60698218067
```

Predict results

```
predictions = model.predict(X_test)  
  
np.set_printoptions(threshold=5) # Print only 5 elements per array  
print(predictions)  
  
[112866.76155431 308640.3343831 120052.08531031 ... 193394.87353075  
 122311.0076527 82585.36875711]  
  
sns.scatterplot(x=y_test, y=predictions)  
plt.xlabel('Actual Values')  
plt.ylabel('Predictions')  
plt.title('Actual prices of houses vs. Model Predictions')  
plt.show()
```



Evaluation of the model


```
from sklearn.metrics import mean_squared_error, mean_absolute_error
import math

print('Mean Absolute Error:', mean_absolute_error(y_test, predictions))
print('Mean Squared Error:', mean_squared_error(y_test, predictions))
print('Root Mean Squared Error:', math.sqrt(mean_squared_error(y_test,
predictions)))
```

```
Mean Absolute Error: 31410.205502278408
Mean Squared Error: 2269069226.157571
Root Mean Squared Error: 47634.74809587609
```


```
residuals = y_test - predictions
sns.histplot(residuals, bins=30)
```

```
<Axes: xlabel='SalePrice', ylabel='Count'>
```




```
import pylab
import scipy.stats as stats

stats.probplot(residuals, dist="norm", plot=pylab)
pylab.show()
```



Use the trained model on our test data

```
Xt = df_test_cleaned[features]
test_predictions = model.predict(Xt)
test_predictions
array([104689.28980193, 159422.08883339, 192805.94429241, ...,
       128229.8257631 , 101461.04973747, 220250.83345123])
plt.title('Salary Distribution Plot')
sns.histplot(test_predictions)
plt.show()
```

```
submission_df = pd.DataFrame({  
    'Id': df_test_cleaned['Id'],  
    'SalePrice': test_predictions  
})  
  
# Save the predictions to a CSV file  
submission_df.to_csv('submission.csv', index=False)
```

Conclusion

The current linear regression model provides a basic understanding of house price predictions, but the relatively high error metrics indicate room for improvement.