

# ML Task-01

Implement a linear regression model to predict the prices of houses based on their square footage and the number of bedrooms and bathrooms.

Dataset : - <https://www.kaggle.com/c/house-prices-advanced-regression-techniques/data>  
(<https://www.kaggle.com/c/house-prices-advanced-regression-techniques/data>)

```
In [24]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [25]: df_train = pd.read_csv('train.csv')
df_test = pd.read_csv('test.csv')
```

```
In [26]: df_train.head()
```

Out[26]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	L
0	1	60	RL	65.0	8450	Pave	NaN	Reg		Lvl
1	2	20	RL	80.0	9600	Pave	NaN	Reg		Lvl
2	3	60	RL	68.0	11250	Pave	NaN	IR1		Lvl
3	4	70	RL	60.0	9550	Pave	NaN	IR1		Lvl
4	5	60	RL	84.0	14260	Pave	NaN	IR1		Lvl

5 rows × 81 columns

In [27]: `df_test.head()`

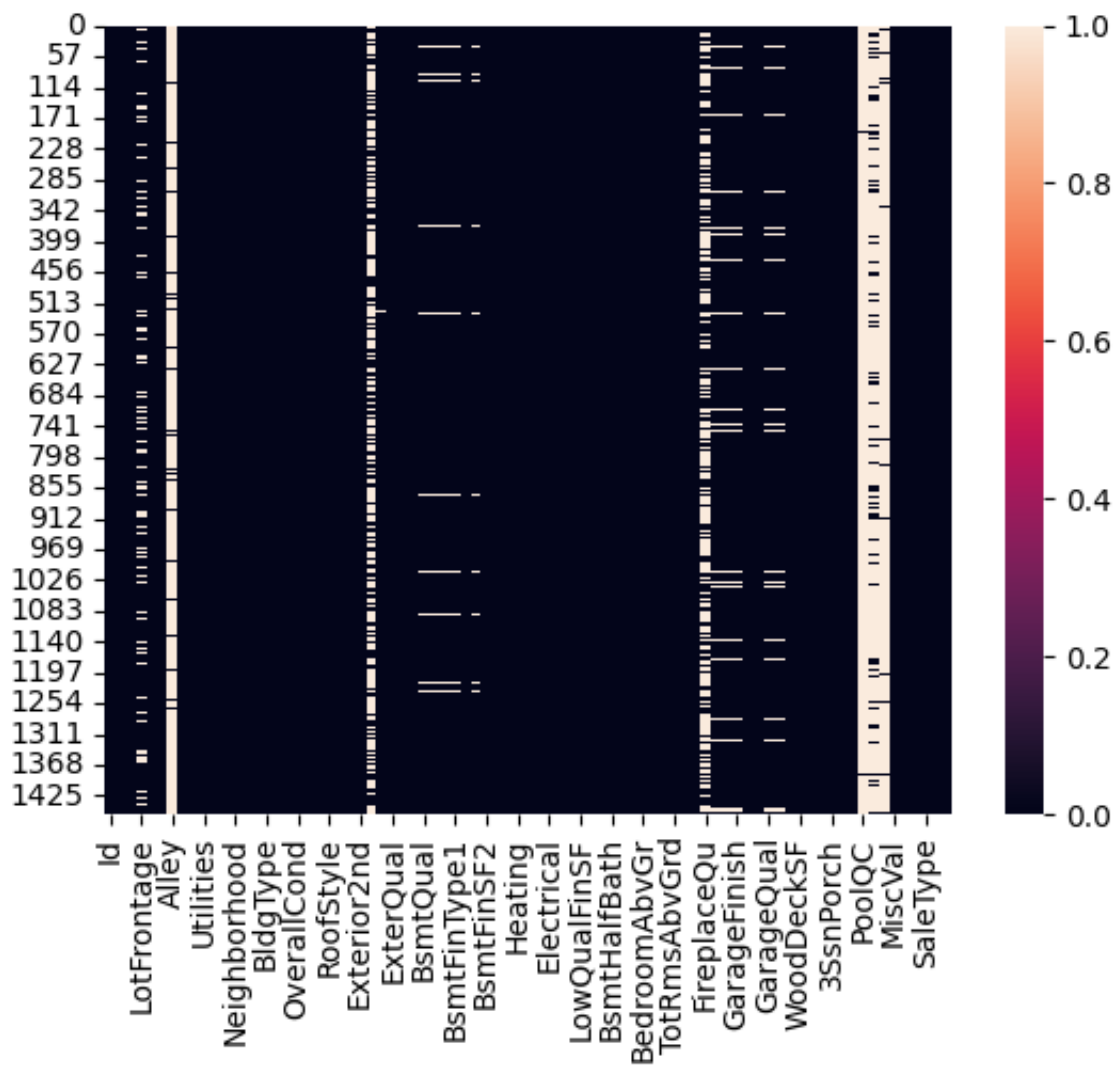
Out[27]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour
0	1461	20	RH	80.0	11622	Pave	NaN	Reg	Lvl
1	1462	20	RL	81.0	14267	Pave	NaN	IR1	Lvl
2	1463	60	RL	74.0	13830	Pave	NaN	IR1	Lvl
3	1464	60	RL	78.0	9978	Pave	NaN	IR1	Lvl
4	1465	120	RL	43.0	5005	Pave	NaN	IR1	HLS

5 rows × 10 columns

```
In [28]: sns.heatmap(df_train.isnull())
```

```
Out[28]: <Axes: >
```



In [29]: `df_train.columns`

Out[29]: Index(['Id', 'MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street',  
 'Alley', 'LotShape', 'LandContour', 'Utilities', 'LotConfig',  
 'LandSlope', 'Neighborhood', 'Condition1', 'Condition2', 'BldgType',  
 'HouseStyle', 'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd',  
 'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType',  
 'MasVnrArea', 'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual',  
 'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1',  
 'BsmtFinType2', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', 'Heating',  
 'HeatingQC', 'CentralAir', 'Electrical', '1stFlrSF', '2ndFlrSF',  
 'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath',  
 'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual',  
 'TotRmsAbvGrd', 'Functional', 'Fireplaces', 'FireplaceQu', 'GarageType',  
 'GarageYrBlt', 'GarageFinish', 'GarageCars', 'GarageArea', 'GarageQual',  
 'GarageCond', 'PavedDrive', 'WoodDeckSF', 'OpenPorchSF', 'EnclosedPorch',  
 '3SsnPorch', 'ScreenPorch', 'PoolArea', 'PoolQC',  
 'Fence', 'MiscFeature', 'MiscVal', 'MoSold', 'YrSold', 'SaleType',  
 'SaleCondition', 'SalePrice'],  
 dtype='object')

In [30]: `null_columns=[]  
 for i in df_train.columns.tolist():  
 if df_train[i].isnull().sum() >= 500:  
 null_columns.append(i)  
 print(i,df_train[i].isnull().sum())`

Alley 1369  
 MasVnrType 872  
 FireplaceQu 690  
 PoolQC 1453  
 Fence 1179  
 MiscFeature 1406

In [31]: `df_train = df_train.drop(columns=null_columns)  
 df_test = df_test.drop(columns=null_columns)`

```
In [32]: df_train_cleaned = df_train.ffill()  
df_test_cleaned = df_test.ffill()
```

```
In [33]: df_train_cleaned.head()
```


Out[33]:

	<b>Id</b>	<b>MSSubClass</b>	<b>MSZoning</b>	<b>LotFrontage</b>	<b>LotArea</b>	<b>Street</b>	<b>LotShape</b>	<b>LandContour</b>	<b>Utilities</b>
<b>0</b>	1	60	RL	65.0	8450	Pave	Reg	Lvl	AllPub
<b>1</b>	2	20	RL	80.0	9600	Pave	Reg	Lvl	AllPub
<b>2</b>	3	60	RL	68.0	11250	Pave	IR1	Lvl	AllPub
<b>3</b>	4	70	RL	60.0	9550	Pave	IR1	Lvl	AllPub
<b>4</b>	5	60	RL	84.0	14260	Pave	IR1	Lvl	AllPub

5 rows × 10 columns

```
In [34]: sns.heatmap(df_train_cleaned.isnull())
```

Out[34]: <Axes: >




```
In [35]: df_train_cleaned.describe()
```

```
Out[35]:
```

	Id	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	Y
<b>count</b>	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000	1460
<b>mean</b>	730.500000	56.897260	70.104795	10516.828082	6.099315	5.575342	1971
<b>std</b>	421.610009	42.300571	23.846996	9981.264932	1.382997	1.112799	30
<b>min</b>	1.000000	20.000000	21.000000	1300.000000	1.000000	1.000000	1872
<b>25%</b>	365.750000	20.000000	59.000000	7553.500000	5.000000	5.000000	1954
<b>50%</b>	730.500000	50.000000	70.000000	9478.500000	6.000000	5.000000	1973
<b>75%</b>	1095.250000	70.000000	80.000000	11601.500000	7.000000	6.000000	2000
<b>max</b>	1460.000000	190.000000	313.000000	215245.000000	10.000000	9.000000	2010

8 rows × 38 columns

```
In [36]: plt.title('Salary Distribution Plot')
sns.histplot(df_train_cleaned['SalePrice'])
plt.show()
```



## Split data

```
In [37]: from sklearn.model_selection import train_test_split
```

```
In [38]: features = ['GrLivArea', '1stFlrSF', '2ndFlrSF', 'LowQualFinSF', 'Bsmt  
FullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath', 'BedroomAbvGr']  
  
target = 'SalePrice'
```


```
In [39]: X = df_train_cleaned[features]  
y = df_train_cleaned[target]
```



```
In [40]: df = pd.concat([X, y], axis=1)


# Plotting grid of scatter plots for each feature against the target variable
sns.pairplot(df, x_vars=features, y_vars=[target], kind='scatter',
              plot_kws={'alpha':0.4}, diag_kws={'alpha':0.55, 'bins':4
0})

plt.show()
```




```
In [41]: sns.lmplot(x='GrLivArea',
                    y='SalePrice',
                    data=df_train_cleaned,
                    scatter_kws={'alpha':0.3})
```

Out[41]: <seaborn.axisgrid.FacetGrid at 0x17727b890>




```
In [42]: sns.lmplot(x='BedroomAbvGr',  
                  y='SalePrice',  
                  data=df_train_cleaned,  
                  scatter_kws={'alpha':0.3})
```

Out[42]: <seaborn.axisgrid.FacetGrid at 0x177079670>



```
In [43]: sns.lmplot(x='FullBath',  
                  y='SalePrice',  
                  data=df_train_cleaned,  
                  scatter_kws={'alpha':0.3})
```

Out[43]: <seaborn.axisgrid.FacetGrid at 0x176f73ec0>



```
In [44]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =  
                  0.2, random_state = 42)
```

## Train model (Linear Regression)

```
In [45]: from sklearn.linear_model import LinearRegression  
  
model = LinearRegression()
```

```
In [46]: model.fit(X_train, y_train)
```

```
Out[46]: 

LinearRegression ⓘ ?  
(https://scikit-learn.org/1.4/modules/generated/sklearn.linear_model.LinearRegression())


```

```
In [47]: coefficients = model.coef_
```

```
In [48]: model.score(X, y)
```

```
Out[48]: 0.6541686106425739
```

```
In [49]: for feature, coef in zip(features, coefficients):  
         print(f'{feature}: {coef}')
```

```
GrLivArea: 45.45358200866404  
1stFlrSF: 70.93123244880826  
2ndFlrSF: 19.004245633922277  
LowQualFinSF: -44.481896073804585  
BsmtFullBath: 22125.53078348177  
BsmtHalfBath: 6538.890647874652  
FullBath: 34869.22887902238  
HalfBath: 22717.7813535538  
BedroomAbvGr: -18379.60698218067
```


## Predict results

```
In [50]: predictions = model.predict(X_test)
```

```
In [51]: np.set_printoptions(threshold=5) # Print only 5 elements per array  
print(predictions)
```

```
[112866.76155431 308640.3343831 120052.08531031 ... 193394.87353075  
122311.0076527 82585.36875711]
```

```
In [52]: sns.scatterplot(x=y_test, y=predictions)
plt.xlabel('Actual Values')
plt.ylabel('Predictions')
plt.title('Actual prices of houses vs. Model Predictions')
plt.show()
```



## Evaluation of the model


```
In [53]: from sklearn.metrics import mean_squared_error, mean_absolute_error
import math
```

```
In [54]: print('Mean Absolute Error:',mean_absolute_error(y_test, predictions))
print('Mean Squared Error:',mean_squared_error(y_test, predictions))
print('Root Mean Squared Error:',math.sqrt(mean_squared_error(y_test,
predictions)))
```

```
Mean Absolute Error: 31410.205502278408
Mean Squared Error: 2269069226.157571
Root Mean Squared Error: 47634.74809587609
```


```
In [55]: residuals = y_test-predictions
sns.histplot(residuals, bins=30)
```

```
Out[55]: <Axes: xlabel='SalePrice', ylabel='Count'>
```



```
In [56]: import pylab
import scipy.stats as stats

stats.probplot(residuals, dist="norm", plot=pylab)
pylab.show()
```



## Use the trained model on our test data

```
In [57]: Xt = df_test_cleaned[features]
```


```
In [58]: test_predictions = model.predict(Xt)
```

```
In [59]: test_predictions
```

```
Out[59]: array([104689.28980193, 159422.08883339, 192805.94429241, ...,
               128229.8257631 , 101461.04973747, 220250.83345123])
```



```
In [60]: plt.title('Salary Distribution Plot')
sns.histplot(test_predictions)
plt.show()
```



```
In [61]: submission_df = pd.DataFrame({
    'Id': df_test_cleaned['Id'],
    'SalePrice': test_predictions
})

# Save the predictions to a CSV file
submission_df.to_csv('submission.csv', index=False)
```

## Conclusion

The current linear regression model provides a basic understanding of house price predictions, but the relatively high error metrics indicate room for improvement.